

 Generate

print hello world using rot13



Close

```
import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from torch.utils.data import DataLoader, SubsetRandomSampler
from torchvision import datasets, transforms
import time

BATCH_SIZE = 64
NUM_EPOCHS = 20
DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])
train_dataset = datasets.MNIST(root='./data',
    train=True,
    download=True,
    transform=transform)
valid_dataset = datasets.MNIST(root='./data',
    train=True,
    transform=transform)
test_dataset = datasets.MNIST(root='./data',
    train=False,
    transform=transform)

validation_fraction = 0.1
num = int(validation_fraction * 60000)
train_indices = torch.arange(0, 60000 - num)
valid_indices = torch.arange(60000 - num, 60000)

train_sampler = SubsetRandomSampler(train_indices)
valid_sampler = SubsetRandomSampler(valid_indices)

train_loader = DataLoader(dataset=train_dataset,
    batch_size=BATCH_SIZE,
    drop_last=True,
    sampler=train_sampler)
valid_loader = DataLoader(dataset=valid_dataset,
    batch_size=BATCH_SIZE,
    sampler=valid_sampler)
test_loader = DataLoader(dataset=test_dataset,
    batch_size=BATCH_SIZE,
    shuffle=False)

# Checking the dataset
for images, labels in train_loader:
    print('Image batch dimensions:', images.shape)
    print('Image label dimensions:', labels.shape)
    break

    Image batch dimensions: torch.Size([64, 1, 28, 28])
    Image label dimensions: torch.Size([64])

import matplotlib.pyplot as plt
# Display a grid of sample images
plt.figure(figsize=(10, 10))
for i, (images, labels) in enumerate(train_loader):
    for j in range(25):
        plt.subplot(5, 5, j + 1)
        plt.imshow(images[j].squeeze(), cmap='gray')
        plt.axis('off')
    break
plt.show()
```



```

# Define your optimizer (using Adam)
optimizer = optim.Adam(model.parameters(), lr=0.001)

import torch.optim as optim
from torch.optim import lr_scheduler
# Create the scheduler with step decay
scheduler = lr_scheduler.StepLR(optimizer, step_size=30, gamma=0.1)

# Create the scheduler with exponential decay
scheduler = lr_scheduler.ExponentialLR(optimizer, gamma=0.95)

def compute_accuracy(data_loader):
    with torch.no_grad():
        correct_pred, num_examples = 0, 0
        for i, (features, targets) in enumerate(data_loader):
            features = features.to(DEVICE)
            targets = targets.float().to(DEVICE)
            logits = model(features)
            _, predicted_labels = torch.max(logits, 1)
            num_examples += targets.size(0)
            correct_pred += (predicted_labels == targets).sum()
        return correct_pred.float()/num_examples * 100

start_time = time.time()
minibatch_loss_list, train_acc_list, valid_acc_list = [], [], []
for epoch in range(NUM_EPOCHS):

    model.train()

    for batch_idx, (features, targets) in enumerate(train_loader):

        features = features.to(DEVICE)
        targets = targets.to(DEVICE)
    # ## FORWARD AND BACK PROP
        logits = model(features)
    #loss = F.cross_entropy(logits, targets)
        loss = criterion(logits, targets)
        optimizer.zero_grad()
        loss.backward()
    # ## UPDATE MODEL PARAMETERS
        optimizer.step()
    # ## LOGGING
        minibatch_loss_list.append(loss.item())
        logging_interval = 100
        if not batch_idx % logging_interval:
            print("Epoch: ", epoch+1, "/", NUM_EPOCHS, "| Batch ", batch_idx,
                  "/", len(train_loader), f'| Loss: {loss:.4f}')
            model.eval()
            with torch.no_grad():# save memory during inference

                train_acc = compute_accuracy(train_loader)
                valid_acc = compute_accuracy(valid_loader)
                print("Epoch: ", epoch+1, "/", NUM_EPOCHS,
                      f'| Train: {train_acc :.2f}% '
                      f'| Validation: {valid_acc :.2f}%')
                train_acc_list.append(train_acc.item())
                valid_acc_list.append(valid_acc.item())
                elapsed = (time.time() - start_time)/60
                print("Time elapsed: ", elapsed, " min")
                scheduler.step(minibatch_loss_list[-1])
                elapsed = (time.time() - start_time)/60
                print(f'Total Training Time: {elapsed:.2f} min')
                test_acc = compute_accuracy(test_loader)
                print(f'Test accuracy {test_acc :.2f}%')

```

```

total training time: 58.74 min
Test accuracy 97.78%
Epoch: 20 / 20 | Batch 100 / 843 | Loss: 0.0402
Epoch: 20 / 20 | Train: 99.29% | Validation: 98.13%
Time elapsed: 59.08817561070124 min
Total Training Time: 59.09 min
Test accuracy 97.66%
Epoch: 20 / 20 | Batch 200 / 843 | Loss: 0.0071
Epoch: 20 / 20 | Train: 99.30% | Validation: 98.12%
Time elapsed: 59.415673883756 min
Total Training Time: 59.42 min
Test accuracy 97.77%
Epoch: 20 / 20 | Batch 300 / 843 | Loss: 0.0114
Epoch: 20 / 20 | Train: 99.51% | Validation: 98.10%
Time elapsed: 59.76093515952428 min
Total Training Time: 59.76 min
Test accuracy 97.82%
Epoch: 20 / 20 | Batch 400 / 843 | Loss: 0.0095
Epoch: 20 / 20 | Train: 99.31% | Validation: 97.77%
Time elapsed: 60.11885837713877 min
Total Training Time: 60.12 min
Test accuracy 97.63%
Epoch: 20 / 20 | Batch 500 / 843 | Loss: 0.0156
Epoch: 20 / 20 | Train: 99.31% | Validation: 97.77%
Time elapsed: 60.45131976207097 min
Total Training Time: 60.45 min
Test accuracy 97.72%
Epoch: 20 / 20 | Batch 600 / 843 | Loss: 0.0118
Epoch: 20 / 20 | Train: 99.19% | Validation: 97.85%
Time elapsed: 60.7944786588351 min
Total Training Time: 60.79 min
Test accuracy 97.56%
Epoch: 20 / 20 | Batch 700 / 843 | Loss: 0.0392
Epoch: 20 / 20 | Train: 99.47% | Validation: 98.08%
Time elapsed: 61.13306200106938 min
Total Training Time: 61.13 min
Test accuracy 97.92%
Epoch: 20 / 20 | Batch 800 / 843 | Loss: 0.0284
Epoch: 20 / 20 | Train: 99.43% | Validation: 97.95%
Time elapsed: 61.4774317463239 min
Total Training Time: 61.48 min
Test accuracy 97.82%


```

```

import matplotlib.pyplot as plt
plt.plot(range(len(minibatch_loss_list)), minibatch_loss_list)
plt.xlabel('Iterations')
plt.ylabel('Loss')

```

```
Text(0, 0.5, 'Loss')
```



```

import numpy as np
num_epochs = len(train_acc_list)
plt.plot(np.arange(1, num_epochs+1),
train_acc_list, label='Training')
plt.plot(np.arange(1, num_epochs+1),
valid_acc_list, label='Validation')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

```

<matplotlib.legend.Legend at 0x7fa516a962c0>

